

Diffusion-LM

Improves **Controllable** Text Generation



NEURAL INFORMATION
PROCESSING SYSTEMS 2022

Xiang Lisa Li, John Thickstun, Ishaan Gulrajani,
Percy Liang, Tatsunori Hashimoto

Stanford University

Pretrained AR-LM

Pre-trained large Autoregressive Language Models (AR-LMs) can generate high-quality text, but in order to be applicable to daily life, it is **necessary to control** the generated text.

It is common to use pairs of (control, text) to fine-tune the model

Issue


- **Updating LM** parameters for different tasks is expensive.
- Difficult to **combine** multiple conditions.
- The **fixed generation order** limits the models' flexibility in many controllable generation settings

Controllable

Plug & Play

$$p(x_{i+1}|x_{0:i}, a) \\ \propto p(a|x_{0:i+1})p(x_{i+1}|x_{0:i})$$

- A pretrained large LM is responsible for generating fluent text.
- Use additional small models to control the attributes of the generated text.
- No need to fine-tune the parameters of the LM.
- Only controls at the attribute level are available.



Diffusion probabilistic models (DPMs) have demonstrated great success in **continuous data domains**, producing images and audio that have state-of-the-art sample quality.

And can use gradient-based control methods for effective control. However...

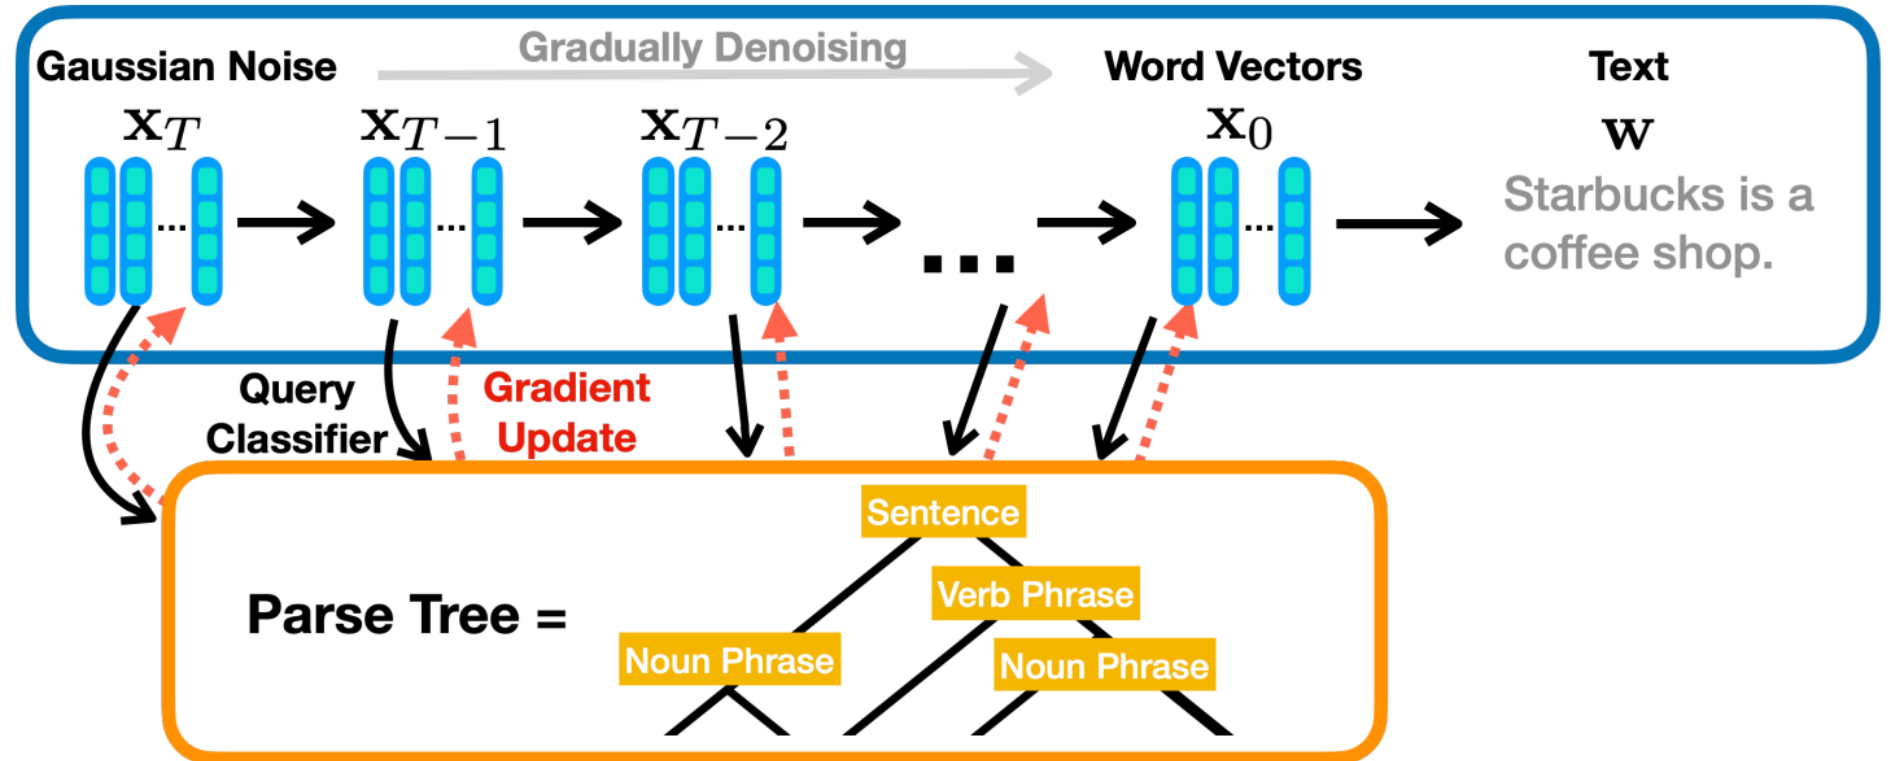
Past Works in Text

- **Discrete** DPMs.
- Unable to use gradient-based control methods.

Diffusion-LM

Continuous DPMs for Text

Diffusion-LM




Classifier

Diffusion-LM

Continuous DPMs for Text

Diffusion-LM




- Starts with a sequence of Gaussian noise vectors
- Incrementally denoises them into vectors corresponding to words.

$$x_T \sim \mathcal{N}(0, I) \in \mathbb{R}^{L \times d}$$

Diffusion-LM

Continuous DPMs for Text

Diffusion-LM

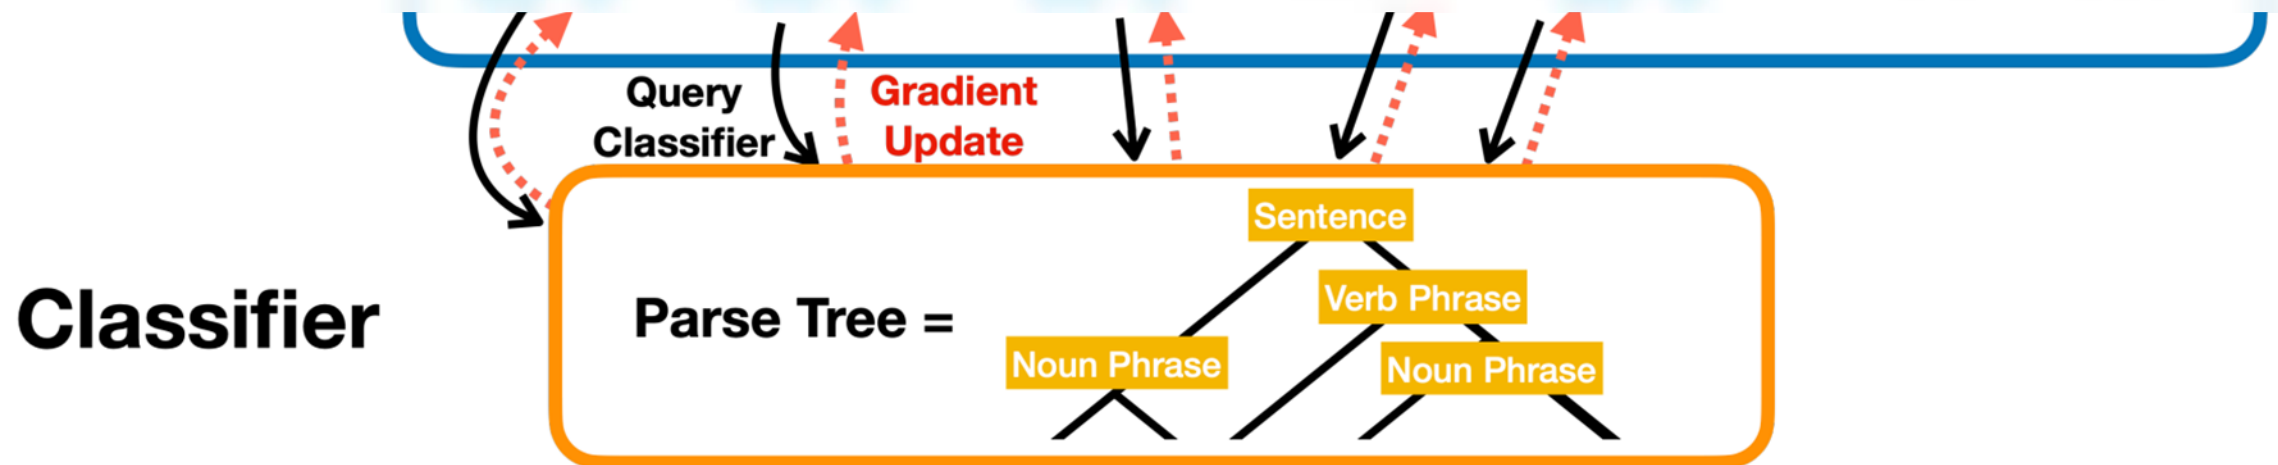


- Starts with a sequence of Gaussian noise vectors
- Incrementally denoises them into vectors corresponding to words.
- There is no fixed generation order so it can be used directly for infilling tasks.

Diffusion-LM

Gradient-based Control

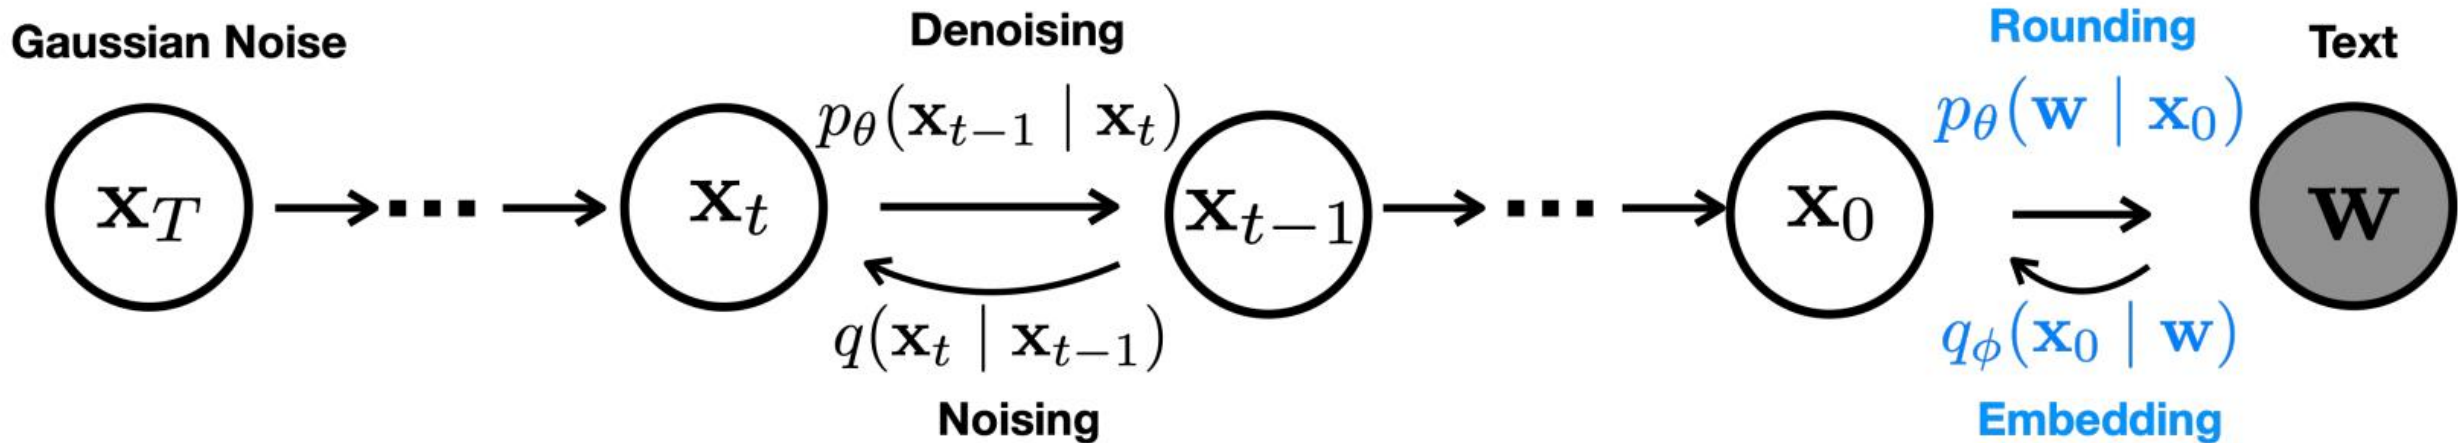
- Change sampling paths using Gradient-based Controls common to DPMs.
- Tested on six tasks, with control over text content and grammatical structure, to demonstrate the controllability of Diffusion-LM.



Diffusion-LM

Loss (MSE ver.)


$$\mathcal{L}_{\text{simple}}^{\text{e2e}}(w) = [\underbrace{\|\mu_{\theta}(\mathbf{x}_t, t) - \hat{\mu}(\mathbf{x}_t, \mathbf{x}_0)\|^2}_{\text{MSE}} - \underbrace{\log p_{\theta}(w|\mathbf{x}_0)}_{\text{Cross Entropy}}]$$



Diffusion-LM

Loss (MSE ver.)

$$\mathcal{L}_{\text{simple}}^{\text{e2e}}(w) = [\underbrace{\|\mu_{\theta}(\mathbf{x}_t, t) - \hat{\mu}(\mathbf{x}_t, \mathbf{x}_0)\|^2}_{\text{MSE}} - \underbrace{\log p_{\theta}(w | \mathbf{x}_0)}_{\text{Cross Entropy}}]$$




$$\hat{\mu}(\mathbf{x}_t, \mathbf{x}_0) = \mathbf{x}_{t-1}$$

Diffusion-LM

Rounding Errors

Empirically, the model fails to generate \mathbf{x}_0 that commits to a single word.

Only starting the denoising at t close to 0 can make \mathbf{x}_0 converge to the word embedding.




Diffusion-LM

Reducing Rounding Errors

This objective quickly learn that \mathbf{x}_0 should precisely centered at a word embedding.

$$\mathcal{L}_{\text{simple}}^{\text{e2e}}(w) = [\underbrace{\|\mu_{\theta}(\mathbf{x}_t, t) - \mathbf{x}_0\|^2}_{\text{Denoising}} - \underbrace{\log p_{\theta}(w | \mathbf{x}_0)}_{\text{Rounding}}]$$



Sampling Algorithm

θ : parameter of Diffusion-LM, φ : parameter of Attribute Model

$\mathbf{x}_T \sim \mathcal{N}(0, I) \in \mathbb{R}^{L \times d}$

for t in $\{T, \dots, 1\}$:

$\mu_t \leftarrow \sqrt{\bar{a}_t} \cdot \text{Clamp}(\mu_\theta(\mathbf{x}_t, t))$

$\mathbf{x}_{t-1} \sim \mathcal{N}(\mu_t, 1 - \bar{a}_t) \in \mathbb{R}^{L \times d}$

for k in $\{1, \dots, K\}$:

$\Delta \mathbf{x}_{t-1} \leftarrow \nabla_{\mathbf{x}_{t-1}} [\lambda \|\mu_t - \mathbf{x}_{t-1}\|^2 - \log p_\varphi(c|\mathbf{x}_{t-1})]$

$\mathbf{x}_{t-1} \leftarrow \text{Adagrad}(\mathbf{x}_{t-1}, \Delta \mathbf{x}_{t-1})$

return $p_\theta(w|\mathbf{x}_0)$

- **Replace** the predicted \mathbf{x}_0 with the **closest word embedding**.
- Applying the clamping trick to early diffusion steps with **t near T may be sub-optimal**.

Controllable Generation

Multiple Gradient Steps

Sampling Algorithm

θ : parameter of Diffusion-LM, φ : parameter of Attribute Model

$\mathbf{x}_T \sim \mathcal{N}(0, I) \in \mathbb{R}^{L \times d}$

for t in $\{T, \dots, 1\}$:

$\mu_t \leftarrow \sqrt{\bar{a}_t} \cdot \text{Clamp}(\mu_\theta(\mathbf{x}_t, t))$

$\mathbf{x}_{t-1} \sim \mathcal{N}(\mu_t, 1 - \bar{a}_t) \in \mathbb{R}^{L \times d}$

for k in $\{1, \dots, K\}$:

$\Delta \mathbf{x}_{t-1} \leftarrow \nabla_{\mathbf{x}_{t-1}} [\lambda \|\mu_t - \mathbf{x}_{t-1}\|^2 - \log p_\varphi(c|\mathbf{x}_{t-1})]$

$\mathbf{x}_{t-1} \leftarrow \text{Adagrad}(\mathbf{x}_{t-1}, \Delta \mathbf{x}_{t-1})$

return $p_\theta(w|\mathbf{x}_0)$

Increase control strength

Sampling Algorithm

θ : parameter of Diffusion-LM, φ : parameter of Attribute Model

$\mathbf{x}_T \sim \mathcal{N}(0, I) \in \mathbb{R}^{L \times d}$

for t in $\{T, \dots, 1\}$:

$\mu_t \leftarrow \sqrt{\bar{a}_t} \cdot \text{Clamp}(\mu_\theta(\mathbf{x}_t, t))$

$\mathbf{x}_{t-1} \sim \mathcal{N}(\mu_t, 1 - \bar{a}_t) \in \mathbb{R}^{L \times d}$

for k in $\{1, \dots, K\}$:

$\Delta \mathbf{x}_{t-1} \leftarrow \nabla_{\mathbf{x}_{t-1}} [\lambda \|\mu_t - \mathbf{x}_{t-1}\|^2 - \log p_\varphi(c|\mathbf{x}_{t-1})]$

$\mathbf{x}_{t-1} \leftarrow \text{Adagrad}(\mathbf{x}_{t-1}, \Delta \mathbf{x}_{t-1})$

return $p_\theta(w|\mathbf{x}_0)$

maintain fluency

Sampling Algorithm

θ : parameter of Diffusion-LM, φ : parameter of Attribute Model

$\mathbf{x}_T \sim \mathcal{N}(0, I) \in \mathbb{R}^{L \times d}$

for t in $\{T, \dots, 1\}$:

$\mu_t \leftarrow \sqrt{\bar{a}_t} \cdot \text{Clamp}(\mu_\theta(\mathbf{x}_t, t))$

$\mathbf{x}_{t-1} \sim \mathcal{N}(\mu_t, 1 - \bar{a}_t) \in \mathbb{R}^{L \times d}$

for k in $\{1, \dots, K\}$:

$\Delta \mathbf{x}_{t-1} \leftarrow \nabla_{\mathbf{x}_{t-1}} [\lambda \|\mu_t - \mathbf{x}_{t-1}\|^2 - \log p_\varphi(c|\mathbf{x}_{t-1})]$

$\mathbf{x}_{t-1} \leftarrow \text{Adagrad}(\mathbf{x}_{t-1}, \Delta \mathbf{x}_{t-1})$

return $p_\theta(w|\mathbf{x}_0)$

Decoding by Minimum Bayes Risk

Control Tasks

Examples

input (Semantic Content)	food : Japanese
output text	Browns Cambridge is good for Japanese food and also children friendly near The Sorrento .
input (Parts-of-speech)	PROPN AUX DET ADJ NOUN NOUN VERB ADP DET NOUN ADP DET NOUN PUNCT
output text	Zizzi is a local coffee shop located on the outskirts of the city .
input (Syntax Tree)	(TOP (S (NP (*) (*) (*)) (VP (*) (NP (NP (*) (*))))))
output text	The Twenty Two has great food
input (Syntax Spans)	(7, 10, VP)
output text	Wildwood pub serves multicultural dishes and is ranked 3 stars
input (Length)	14
output text	Browns Cambridge offers Japanese food located near The Sorrento in the city centre .
input (left context)	My dog loved tennis balls.
input (right context)	My dog had stolen every one and put it under there.
output text	One day, I found all of my lost tennis balls underneath the bed.

Training Dataset. E2E & ROCStories. Challenge. ROCStories > E2E.

Control Tasks

Semantic Content

input (Semantic Content)	food : Japanese
output text	Browns Cambridge is good for Japanese food and also children friendly near The Sorrento .

Given a **field** and **value**, generate a sentence that covers field=value.

- Classifier-Based
- Train an autoregressive LM (GPT-2 small architecture) to predict the (field, value) pair.
 $\langle \text{value} \rangle \leftarrow \text{AR-LM}(\langle \text{sentence} \rangle \langle \text{eos} \rangle \langle \text{field} \rangle \langle \text{eos} \rangle)$

Control Tasks

Parts-of-speech

Generate a sequence of words of the same length whose POS tags match the target.

input (Parts-of-speech)	PROPN AUX DET ADJ NOUN NOUN VERB ADP DET NOUN ADP DET NOUN PUNCT
output text	Zizzi is a local coffee shop located on the outskirts of the city .

- Classifier-Based
- BERT-base architecture

parametrized by a parts-of-speech tagger, which estimates the probability of the target POS sequence conditioned on the latent variables.

Given a target syntactic parse tree, generate text whose syntactic parse matches the given parse.

input (Syntax Tree)	(TOP (S (NP (*) (*) (*)) (VP (*) (NP (NP (*) (*))))))
output text	The Twenty Two has great food

- Classifier-Based
- Transformer-based constituency parser
Constituency parsing with a self-attentive encoder, ACL, 2018

Control Tasks

Syntax Spans

Given a target (span, syntactic category) pair, generate text whose parse tree over span $[i, j]$ matches the target syntactic category.

input (Syntax Spans)	(7, 10, VP)
output text	Wildwood pub serves multicultural dishes and is ranked 3 stars

- Classifier-based
- Use the same parser trained for the syntax tree

Control Tasks

Length & Infilling

Length. Given a target length and generate a sequence with a length within ± 2 of the target.

Infilling. Given a left context and a right context, and generate a sentence to connect them.

- Classifier-Free

input (Length)	14
output text	Browns Cambridge offers Japanese food located near The Sorrento in the city centre .
input (left context)	My dog loved tennis balls.
input (right context)	My dog had stolen every one and put it under there.
output text	One day, I found all of my lost tennis balls underneath the bed.

Algorithm 1 Infilling Algorithm

prefix, suffix

infix length

T : number of diffusion steps

p_θ : trained diffusion model

q_θ : forward diffusion process

1: $x_T(\text{infix}) \sim \text{Gaussian}(0, \mathbf{I})$

2: **for** $t \in \{T - 1, \dots, 0\}$ **do**

\mathbb{B} :

$$x_{t+1}(\text{prefix}) \sim q_\theta(x_{t+1} \mid \text{EMB}(\text{prefix}))$$

$$x_{t+1}(\text{suffix}) \sim q_\theta(x_{t+1} \mid \text{EMB}(\text{suffix}))$$

$$x_{t+1} = [x_{t+1}(\text{prefix}), x_{t+1}(\text{infix}), x_{t+1}(\text{suffix})]$$

$$x_t \sim p_\theta(x_t \mid x_{t+1})$$

4:

5: **return** Round($x_0(\text{infix})$)

Main Results

Classifier-Guided

	Semantic Content		Parts-of-speech		Syntax Tree		Syntax Spans		Length	
	ctrl ↑	lm ↓	ctrl ↑	lm ↓	ctrl ↑	lm ↓	ctrl ↑	lm ↓	ctrl ↑	lm ↓
PPLM	9.9	5.32	-	-	-	-	-	-	-	-
FUDGE	69.9	2.83	27.0	7.96	17.9	3.39	54.2	4.03	46.9	3.11
Diffusion-LM	81.2	2.55	90.0	5.16	86.0	3.71	93.8	2.53	99.9	2.16
FT-sample	72.5	2.87	89.5	4.72	64.8	5.72	26.3	2.88	98.1	3.84
FT-search	89.9	1.78	93.0	3.31	76.4	3.24	54.4	2.19	100.0	1.83

Both PPLM and FUDGE are plug-and-play controllable generation approaches based on **AR-LMs trained from scratch using the GPT-2 small architecture**.

lm-score. Feed the generated text to a teacher LM (i.e., a carefully fine-tuned GPT-2 model) and report the perplexity of generated text under the teacher LM.

Main Results

Classifier-Guided

	Semantic Content		Parts-of-speech		Syntax Tree		Syntax Spans		Length	
	ctrl ↑	lm ↓	ctrl ↑	lm ↓	ctrl ↑	lm ↓	ctrl ↑	lm ↓	ctrl ↑	lm ↓
PPLM	9.9	5.32	-	-	-	-	-	-	-	-
FUDGE	69.9	2.83	27.0	7.96	17.9	3.39	54.2	4.03	46.9	3.11
Diffusion-LM	81.2	2.55	90.0	5.16	86.0	3.71	93.8	2.53	99.9	2.16
FT-sample	72.5	2.87	89.5	4.72	64.8	5.72	26.3	2.88	98.1	3.84
FT-search	89.9	1.78	93.0	3.31	76.4	3.24	54.4	2.19	100.0	1.83

	Semantic Content + Syntax Tree			Semantic Content + Parts-of-speech		
	semantic ctrl ↑	syntax ctrl ↑	lm ↓	semantic ctrl ↑	POS ctrl ↑	lm ↓
FUDGE	61.7	15.4	3.52	64.5	24.1	3.52
Diffusion-LM	69.8	74.8	5.92	63.7	69.1	3.46
FT-PoE	61.7	29.2	2.77	29.4	10.5	2.97

lm-score. Feed the generated text to a teacher LM (i.e., a carefully fine-tuned GPT-2 model) and report the perplexity of generated text under the teacher LM.

Main Results

Infilling

	Automatic Eval				Human Eval
	BLEU-4 \uparrow	ROUGE-L \uparrow	CIDEr \uparrow	BERTScore \uparrow	
Left-only	0.9	16.3	3.5	38.5	n/a
DELOREAN	1.6	19.1	7.9	41.7	n/a
COLD	1.8	19.5	10.7	42.7	n/a
Diffusion	7.1	28.3	30.7	89.0	0.37 ^{+0.03} _{-0.02}
AR	6.7	27.0	26.9	89.0	0.39 ^{+0.02} _{-0.03}

AR-infilling. $x_{i+1} \leftarrow \text{AR-LM}(x_{0:l}; x_{r:\text{end}}; x_{l+1:i}), l < i < r$

Conclusions

This study proposes a **novel text generation model** based on a **Continuous Diffusion Model** and exhibits good performance on six different control tasks.

Advantages

- High flexibility and quality
- Achieved performance similar to or even better than FT-AR without fine-tune.

Drawbacks

- Higher perplexity.
- Decoding is substantially slower.
- Training converges more slowly.

- Doubt

GPT 2 in this study was trained from scratch, and it is unknown whether the amount of training data is sufficient.

- Next Version

DiffuSeq: **Sequence to Sequence** Text Generation with Diffusion Models

- Prediction

In the near future, there will be research institutions developing large-scale pre-trained Diffusion-based Language Models.
(Recorded on 2022/11/07)

Additional Information

PPLM

$$p(x_{i+1}|x_{0:i}, a) \\ \propto p(a|x_{0:i+1})p(x_{i+1}|x_{0:i})$$

$$x_{i+1}, H_{i+1} \leftarrow \text{LM} \left(x_i, H_i + \Delta H_i^{(T)} \right)$$

$$\Delta H_i^{(t+1)} \\ \leftarrow \Delta H_i^{(t)} + \alpha \frac{\nabla_{\Delta H_i} \log p(a|H_i + \Delta H_i^{(t)})}{\left\| \nabla_{\Delta H_i} \log p(a|H_i + \Delta H_i^{(t)}) \right\|^\gamma}$$

