

Neural Processes

ICML 2018, ICML 2018 workshop

Introduction

Introduction

A Neural network (NN) is a parameterized function that approximates a dataset with high accuracy, but **can't easily switch to an unknown function**.

Gaussian process (GP) has the flexibility to infer the shape of new functions at test time based on **prior knowledge**, but GP is **computationally intensive**.


Introduction

Therefore, this study proposes a new meta learning method called Neural Process, which combines the advantages of both.

This method is not only as **computationally efficient** as NNs during training and evaluation, but also as GPs can **effectively utilize prior knowledge** to quickly adapt to newly observed functions.

Conditional/Latent Neural Processes

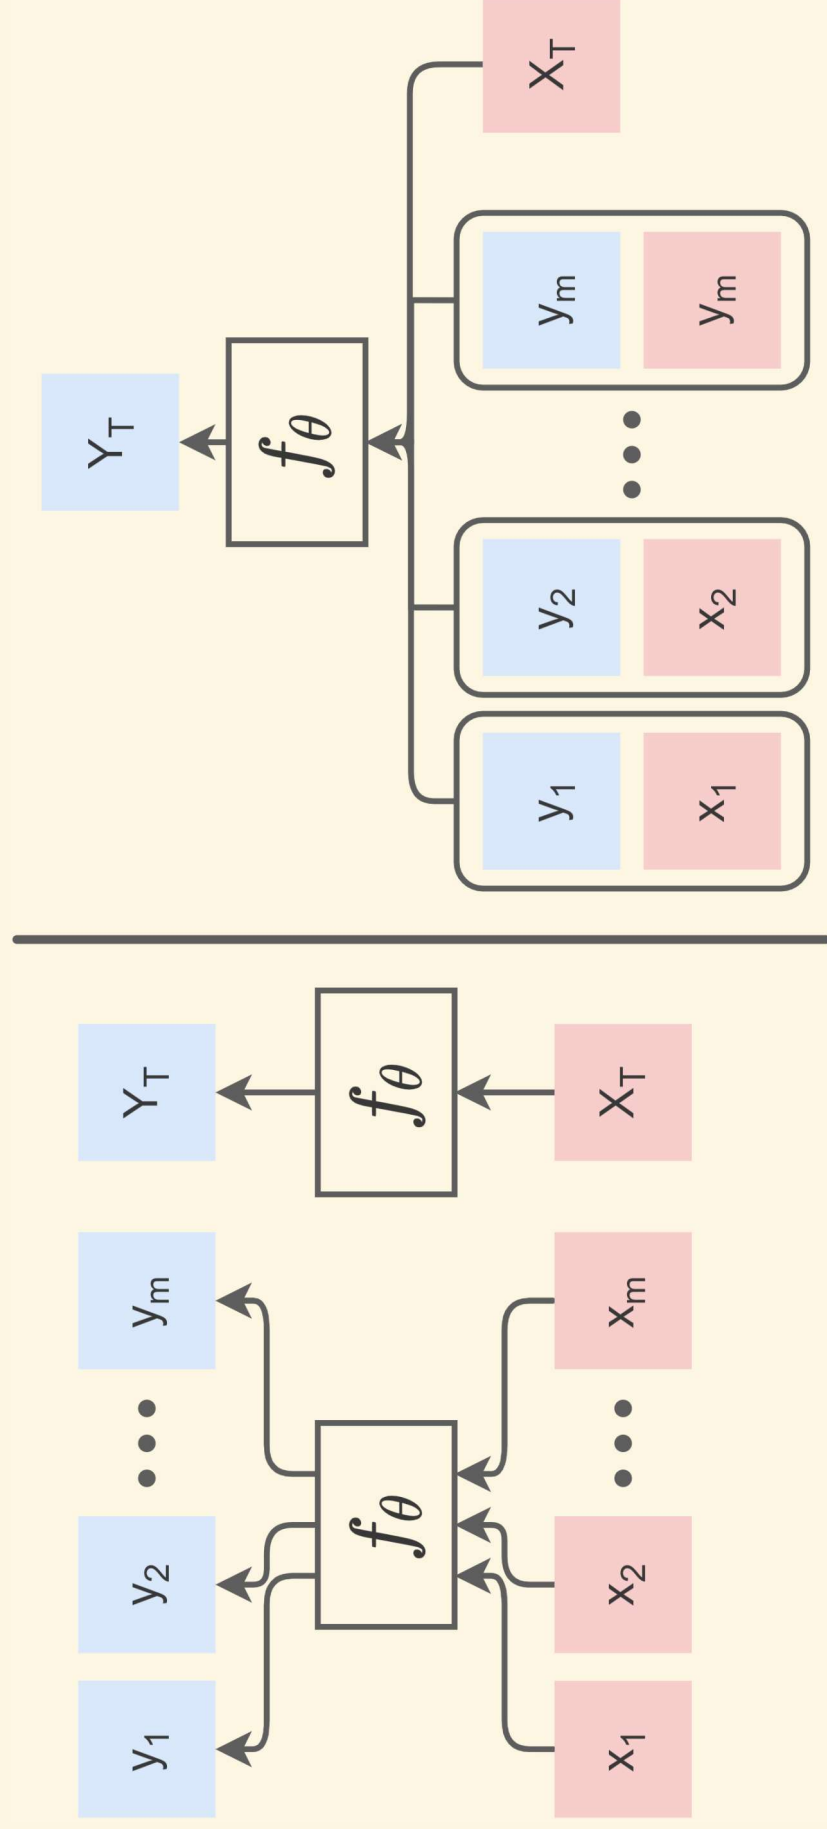
Observation and Target Dataset




Supervised Learning

vs.


Neural Processes




CNP Forward Pass



Conditional Neural Processes



Conditional Neural Processes




$$\mu_{Y_T}, \sigma_{Y_T} \leftarrow Dec(X_T, r)$$


$$\tilde{Y}_T \sim N(\mu_{Y_T}, \sigma_{Y_T})$$

Conditional Neural Processes


$$\log p_{\theta}(Y_T | \{X_O, Y_O\}, X_T)$$



Latent Neural Processes



Latent Neural Processes




$$\mu_O, \sigma_O \leftarrow (MLP \circ mean \circ Enc)(X_O, Y_O)$$

$$z_T \sim N(\mu_O, \sigma_O)$$

Latent Neural Processes

$$\log \int p_{\theta}(Y_T | z, X_T) q_{\omega}(z | \{X_O, Y_O\}) dz$$



Latent Neural Processes

$$-\log p_{\theta}(Y_T|z, X_T) + \log \frac{q_{\omega}(z)}{q_{\omega}(z|\{X_O, Y_O\})}$$


Latent Neural Processes

$$-\log p_{\theta}(Y_T | z, X_T) + \log \frac{q_{\omega}(z)}{q_{\omega}(z | \{X_O, Y_O\})}$$


$$\begin{aligned} L = & -\log P_{\theta}(Y_T | z, X_T) \\ & + KLD(N(\mu_T, \sigma_T), N(\mu_C, \sigma_C)) \end{aligned}$$

Experiments


LNP | CelebA32 | C=0




CNP | CelebA32 | C=0



LNP | MNIST | $C=0$



CNP | MNIST | $C=0$

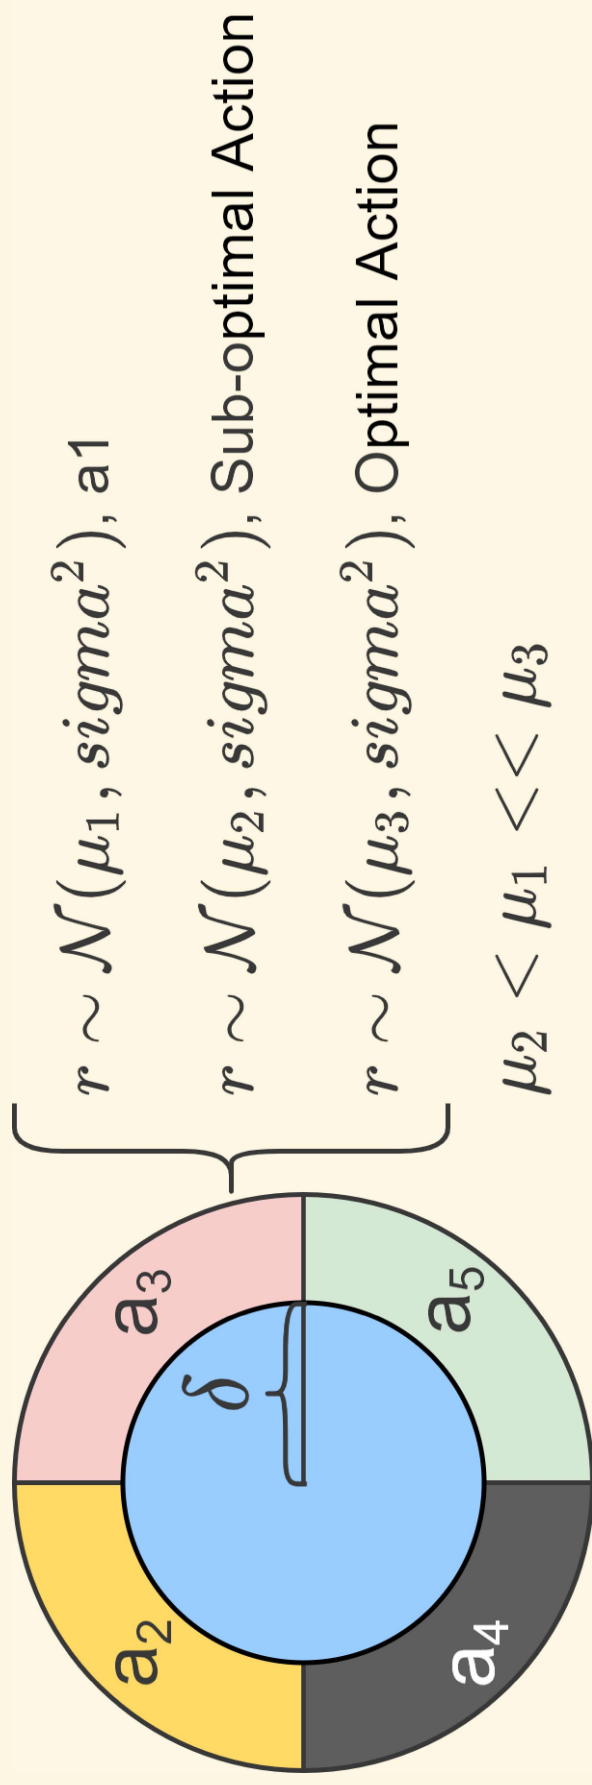


#	Random Context				Ordered Context			
	10	100	1000	1000	10	100	1000	1000
kNN	0.215	0.052	0.007	0.007	0.370	0.273	0.007	0.007
GP	0.247	0.137	0.001		0.257	0.220	0.002	
CNP	0.039	0.016	0.009		0.057	0.047	0.021	

Pixel-wise mean squared error of CNP on the CelebA.

	5-way Acc		20-way Acc		Runtime
	1-shot	5-shot	1-shot	5-shot	
MANN	82.8%	94.9%	-	-	$\mathcal{O}(nm)$
MN	98.1%	98.9%	93.8%	98.5%	$\mathcal{O}(nm)$
CNP	95.3%	98.5%	89.9%	96.8%	$\mathcal{O}(n + m)$

Classification results of CNP on Omniglot.



The wheel bandit problem. Optimal action for yellow, red, black and green regions, are actions 2, 3, 4 and 5, respectively.

δ	0.5	0.7	0.9	0.95	0.99
Cumulative regret					
Uniform	100.00 \pm 0.08	100.00 \pm 0.09	100.00 \pm 0.25	100.00 \pm 0.37	100.00 \pm 0.78
LinGreedy ($\epsilon = 0.0$)	65.89 \pm 4.90	71.71 \pm 4.31	108.86 \pm 3.10	102.80 \pm 3.06	104.80 \pm 0.91
Dropout	7.89 \pm 1.51	9.03 \pm 2.58	36.58 \pm 3.62	63.12 \pm 4.26	98.68 \pm 1.59
LinGreedy ($\epsilon = 0.05$)	7.86 \pm 0.27	9.58 \pm 0.35	19.42 \pm 0.78	33.06 \pm 2.06	74.17 \pm 1.63
Bayes by Backprob (Blundell et al., 2015)	1.37 \pm 0.07	3.32 \pm 0.80	34.42 \pm 5.50	59.04 \pm 5.59	97.38 \pm 2.66
NeuralLinear	0.95 \pm 0.02	1.60 \pm 0.03	4.65 \pm 0.18	9.56 \pm 0.36	49.63 \pm 2.41
MAML (Finn et al., 2017)	2.95 \pm 0.12	3.11 \pm 0.16	4.84 \pm 0.22	7.01 \pm 0.33	22.93 \pm 1.57
Neural Processes	1.60 \pm 0.06	1.75 \pm 0.05	3.31 \pm 0.10	5.71 \pm 0.24	22.13 \pm 1.23
Simple regret					
Uniform	100.00 \pm 0.45	100.00 \pm 0.78	100.00 \pm 1.18	100.00 \pm 2.21	100.00 \pm 4.21
LinGreedy ($\epsilon = 0.0$)	66.59 \pm 5.02	73.06 \pm 4.55	108.56 \pm 3.65	105.01 \pm 3.59	105.19 \pm 4.14
Dropout	6.57 \pm 1.48	6.37 \pm 2.53	35.02 \pm 3.94	59.45 \pm 4.74	102.12 \pm 4.76
LinGreedy ($\epsilon = 0.05$)	5.53 \pm 0.19	6.07 \pm 0.24	8.49 \pm 0.47	12.65 \pm 1.12	57.62 \pm 3.57
Bayes by Backprob (Blundell et al., 2015)	0.60 \pm 0.09	1.45 \pm 0.61	27.03 \pm 6.19	56.64 \pm 6.36	102.96 \pm 5.93
NeuralLinear	0.33 \pm 0.04	0.79 \pm 0.07	2.17 \pm 0.14	4.08 \pm 0.20	35.89 \pm 2.98
MAML (Finn et al., 2017)	2.49 \pm 0.12	3.00 \pm 0.35	4.75 \pm 0.48	7.10 \pm 0.77	22.89 \pm 1.41
Neural Processes	1.04 \pm 0.06	1.26 \pm 0.21	2.90 \pm 0.35	5.45 \pm 0.47	21.45 \pm 1.3

Results of LNP on the wheel bandit problem.

Conclusion

- Neural Processes combines the computational efficiency of neural networks with the flexibility of stochastic processes.
- Efficiently extract prior knowledge from $\{X_O, Y_O\}$
(neural networks)
 - Change strategies based on prior knowledge
(stochastic processes)

Thanks for your attention.

Q&A